

# network

June 6, 2024

```
[ ]: import torch.nn as nn
import torchvision
from torchvision import transforms
import torch
from torch.utils.data import DataLoader
from tqdm import tqdm
import matplotlib.pyplot as plt
import numpy as np
```

## 1 Task 1: Implement a Linear Autoencoder

### 1.0.1 1.1 Create Autoencoder

```
[ ]: # Normalize dataset to the range [-1;1]
transform = transforms.Compose([
    transforms.Resize(size=(28, 28)),
    transforms.ToTensor(), # Range [0;1]
    transforms.Normalize(0.5, 0.5) # Mean and std = 0.5: [0;1] =>
    ↪ [-1;1] since out = (x - mean) / std
])
```

*# Load dataset*

```
dataset_train = torchvision.datasets.FashionMNIST('../data', download=True,
    ↪ train=True, transform=transform)
dataset_test = torchvision.datasets.FashionMNIST('../data', download=True,
    ↪ train=False, transform=transform)
```

```
[ ]: # Visualize some data
image_labels = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
    ↪ 'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']
def plot_examples(data):
    # Plot some examples and put their corresponding label on top as title.
    fig, axs = plt.subplots(1,10, figsize=(20,10))
    for i in range(10):
        img = data.data[i]
        class_id = data.targets[i]
```

```

        axs[i].imshow(img,cmap='gray')
        axs[i].set_title(str(image_labels[class_id.item()]))
        axs[i].axis('off')

fig.show()
# Also print some statistics
print('Statistics:')
print('Min value',int(data.data.min()))
print('Max value',int(data.data.max()))
print('Mean value',float(data.data.float().mean()))
print('Shape',tuple(data.data.shape))
print('Data type',data.data.dtype)

plot_examples(dataset_train)

```


C:\Users\adria\AppData\Local\Temp\ipykernel\_12492\772462545.py:13: UserWarning: FigureCanvasAgg is non-interactive, and thus cannot be shown

```
fig.show()
```

```

Statistics:
Min value 0
Max value 255
Mean value 72.94035339355469
Shape (60000, 28, 28)
Data type torch.uint8

```



```

[ ]: # Define Network
class Autoencoder(nn.Module):

    def __init__(self, input_shape=(28, 28)):
        super(Autoencoder, self).__init__()
        # Compress the feature space of 28*28 to an 8 dimensional
        latent space
        self.encoder = nn.Sequential(
            nn.Linear(in_features=28*28, out_features=128),
            nn.ReLU(),
            nn.Linear(in_features=128, out_features=64),
            nn.ReLU(),
            nn.Linear(in_features=64, out_features=32),
            nn.ReLU(),
            nn.Linear(in_features=32, out_features=16),

```

```

        nn.ReLU(),
        nn.Linear(in_features=16, out_features=8),
        nn.ReLU()
    )

    # Blow up the latent space to the original 28*28 feature space
    self.decoder = nn.Sequential(
        nn.Linear(in_features=8, out_features=8),
        nn.ReLU(),
        nn.Linear(in_features=8, out_features=32),
        nn.ReLU(),
        nn.Linear(in_features=32, out_features=64),
        nn.ReLU(),
        nn.Linear(in_features=64, out_features=128),
        nn.ReLU(),
        nn.Linear(in_features=128, out_features=784),
        nn.Tanh(),
    )

    def forward(self, x):
        # Flatten the 2 dimensional image
        x = torch.flatten(x, start_dim=1, end_dim=-1) # do not flatten
        ↪ batch size

        # Encode image to latent space
        x = self.encoder(x)

        # Reconstruct image
        x = self.decoder(x)

        return x

```

### 1.0.2 1.2 Implement reconstruction loss

```
[ ]: criterion = nn.MSELoss()
```

### 1.0.3 1.3 Perform training + 1.4 Visualize results

```
[ ]: # Hyperparameter
lr = 0.001
weight_decay = 1e-5
batch_size = 128
epochs = 51

# Select the device to work on.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

```

```

model = Autoencoder().to(device)
optimizer = torch.optim.Adam(params=model.parameters(), lr=lr,
    ↳weight_decay=weight_decay)

# Initialize dataloader
dataloader_train = DataLoader(dataset_train, batch_size=batch_size,
    ↳shuffle=True, drop_last=True) # Drop last, to drop last batch since it is
    ↳smaller than 128
dataloader_test = DataLoader(dataset_test, batch_size=batch_size, shuffle=True,
    ↳drop_last=True)

total_loss_train = np.zeros(epochs)
total_loss_test = np.zeros(epochs)
# Perform training
for epoch in range(epochs):
    loss_train = 0
    model.train()
    for [example, _] in dataloader_train:
        example = example.to(device)

        optimizer.zero_grad()
        #print(example.shape)
        prediction = model(example)
        prediction = torch.reshape(prediction, shape=(128, 1, 28, 28))
    ↳# Reconstruct original image dimension
        loss = criterion(prediction, example)

        loss_train += loss

        loss.backward()

        optimizer.step()

    loss_train = loss_train / len(dataloader_train)
    total_loss_train[epoch] = loss_train
    print(f"Epoch: {epoch}, Training loss: {round(loss_train.item(), 4)}")

# Every 10 epochs, also calculate the loss on the test set
if epoch % 10 == 0:
    loss_test = 0

    examples_test = []
    predictions_test = []

    model.eval()
    with torch.no_grad():

```

```

        for [example, _] in dataloader_test:
            example = example.to(device)

            prediction = model(example)
            prediction = torch.reshape(prediction, ↵
↵shape=(128, 1, 28, 28)) # Reconstruct original image dimension
            loss = criterion(prediction, example)

            loss_test += loss
            examples_test.append(example)
            predictions_test.append(prediction)
        loss_test = loss_test / len(dataloader_test)
        total_loss_test[epoch] = loss_test
        print("Test loss: ", round(loss_test.item(), 4))

    fig, axs = plt.subplots(2, 10, figsize=(20, 5))

    # Draw 10 random image indices
    random_nmbrs = np.random.
↵choice(range(len(predictions_test)), size=10)

    for i in range(10):
        example_image = examples_test[i].cpu().
↵detach()[0][0]

        prediction_image = predictions_test[i].cpu().
↵detach()[0][0]

        fig.suptitle(f'Epoch: {epoch}; Training loss: ↵
↵{loss_train:.4f} Testing loss: {loss_test:.4f}')
        axs[0][i].imshow(example_image, cmap='gray')
        axs[0][i].set_title("original")
        axs[0][i].axis('off')
        axs[1][i].imshow(prediction_image, cmap='gray')
        axs[1][i].set_title("reconstruction")
        axs[1][i].axis('off')

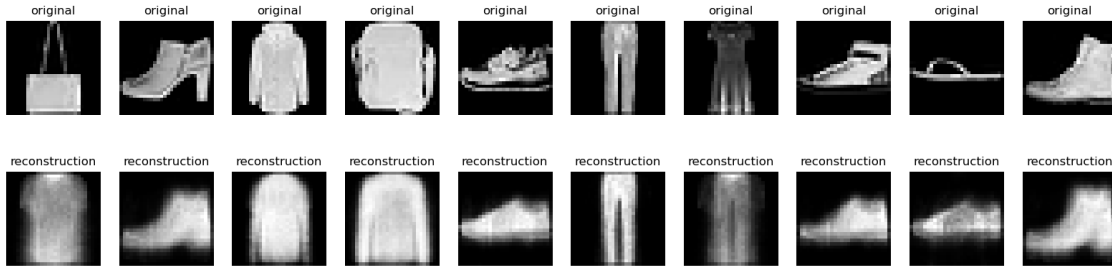
    plt.show()

```

Epoch: 0, Training loss: 0.2008

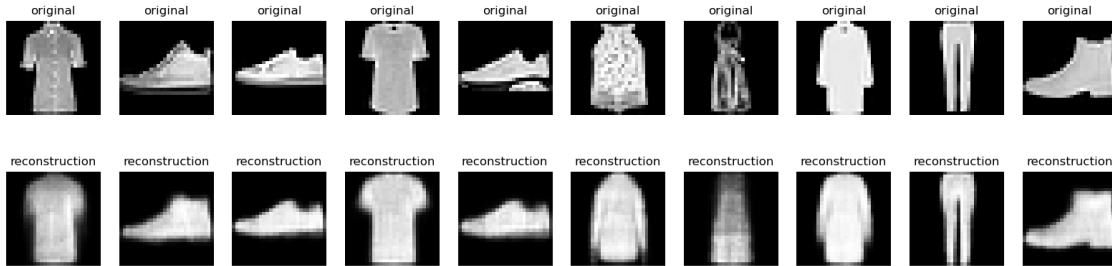
Test loss: 0.1281

Epoch: 0; Training loss: 0.2008 Testing loss: 0.1281



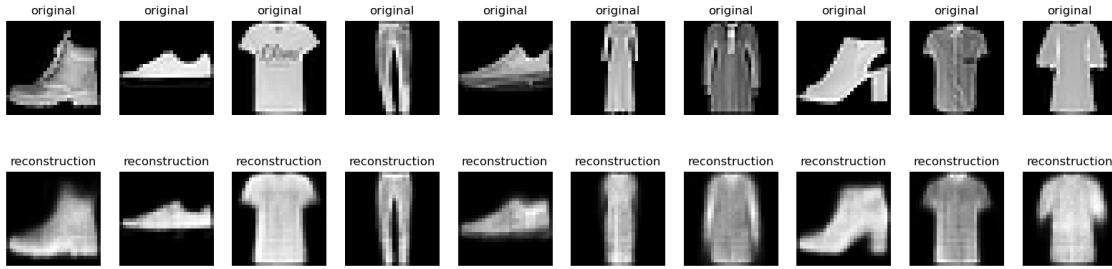
Epoch: 1, Training loss: 0.1192  
Epoch: 2, Training loss: 0.108  
Epoch: 3, Training loss: 0.1029  
Epoch: 4, Training loss: 0.098  
Epoch: 5, Training loss: 0.0923  
Epoch: 6, Training loss: 0.0893  
Epoch: 7, Training loss: 0.0869  
Epoch: 8, Training loss: 0.0847  
Epoch: 9, Training loss: 0.0819  
Epoch: 10, Training loss: 0.0799  
Test loss: 0.0794

Epoch: 10; Training loss: 0.0799 Testing loss: 0.0794



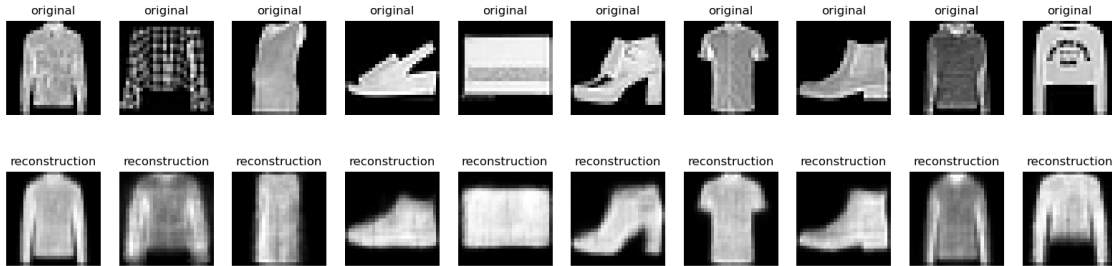
Epoch: 11, Training loss: 0.0788  
Epoch: 12, Training loss: 0.0779  
Epoch: 13, Training loss: 0.0773  
Epoch: 14, Training loss: 0.0766  
Epoch: 15, Training loss: 0.076  
Epoch: 16, Training loss: 0.0755  
Epoch: 17, Training loss: 0.075  
Epoch: 18, Training loss: 0.0745  
Epoch: 19, Training loss: 0.074  
Epoch: 20, Training loss: 0.0737  
Test loss: 0.0751

Epoch: 20; Training loss: 0.0737 Testing loss: 0.0751

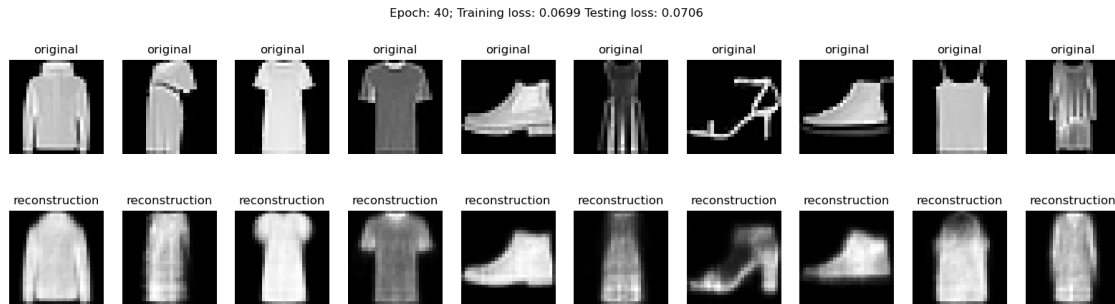


Epoch: 21, Training loss: 0.0734  
Epoch: 22, Training loss: 0.0731  
Epoch: 23, Training loss: 0.0728  
Epoch: 24, Training loss: 0.0726  
Epoch: 25, Training loss: 0.0723  
Epoch: 26, Training loss: 0.0721  
Epoch: 27, Training loss: 0.0719  
Epoch: 28, Training loss: 0.0717  
Epoch: 29, Training loss: 0.0716  
Epoch: 30, Training loss: 0.0714  
Test loss: 0.0724

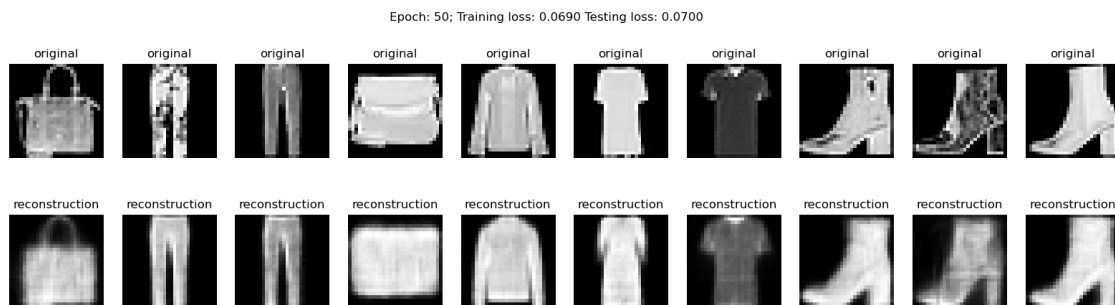
Epoch: 30; Training loss: 0.0714 Testing loss: 0.0724



Epoch: 31, Training loss: 0.0712  
Epoch: 32, Training loss: 0.071  
Epoch: 33, Training loss: 0.0709  
Epoch: 34, Training loss: 0.0707  
Epoch: 35, Training loss: 0.0706  
Epoch: 36, Training loss: 0.0705  
Epoch: 37, Training loss: 0.0704  
Epoch: 38, Training loss: 0.0703  
Epoch: 39, Training loss: 0.0701  
Epoch: 40, Training loss: 0.0699  
Test loss: 0.0706



Epoch: 41, Training loss: 0.0699  
 Epoch: 42, Training loss: 0.0698  
 Epoch: 43, Training loss: 0.0696  
 Epoch: 44, Training loss: 0.0696  
 Epoch: 45, Training loss: 0.0695  
 Epoch: 46, Training loss: 0.0694  
 Epoch: 47, Training loss: 0.0693  
 Epoch: 48, Training loss: 0.0692  
 Epoch: 49, Training loss: 0.0691  
 Epoch: 50, Training loss: 0.069  
 Test loss: 0.07



```
[ ]: # Save the trained model to be used later
torch.save(model.state_dict(), 'autoencoder_{}.ckpt'.format(epoch))
```

## 2 Task 2: Denoising autoencoder

```
[ ]: def add_white_noise(x, factor=0.5, stddev=1):
    """ Adds white noise to an input tensor.
    To make sure that data is in intended range [min, max], use torch.
    ↪ clamp(x, min, max) after applying this function.

    :param x: ND Tensor that is altered
```



```

        :param factor: A factor that controls the strength of the additive noise
        :param stddev: The stddev of the normal distribution used for
        ↪generating the noise
        :return: ND Tensor, x with white noise
        """
        # add white noise to tensor
        noise = x.clone().normal_(0, stddev)
        return x + noise * factor

```

```

[ ]: model = Autoencoder().to(device)
optimizer = torch.optim.Adam(params=model.parameters(), lr=lr,
    ↪weight_decay=weight_decay)

# Initialize dataloader
# Drop last, to drop last batch since it is smaller than 128
dataloader_train = DataLoader(dataset_train, batch_size=batch_size,
    ↪shuffle=True, drop_last=True)
dataloader_test = DataLoader(dataset_test, batch_size=batch_size, shuffle=True,
    ↪drop_last=True)

total_loss_train = np.zeros(epochs)
total_loss_test = np.zeros(epochs)
# Perform training
for epoch in range(epochs):
    loss_train = 0
    model.train()
    for [example, _] in dataloader_train:
        example = example.to(device)

        example_noisy = add_white_noise(example)
        example_noisy = torch.clamp(example_noisy, min=-1, max=1)

        optimizer.zero_grad()
        prediction = model(example_noisy)
        prediction = torch.reshape(prediction, shape=(128, 1, 28, 28))
    ↪# Reconstruct original image dimension
        loss = criterion(prediction, example)

        loss_train += loss

    loss.backward()

    optimizer.step()

    loss_train = loss_train / len(dataloader_train)
    total_loss_train[epoch] = loss_train
    print(f"Epoch: {epoch}, Training loss: {round(loss_train.item(), 4)}")

```

```

# Every 10 epochs, also calculate the loss on the test set
if epoch % 10 == 0:
    loss_test = 0

    examples_test = []
    predictions_test = []

    model.eval()
    with torch.no_grad():
        for [example, _] in dataloader_test:
            example = example.to(device)

            example_noisy = add_white_noise(example)
            example_noisy = torch.clamp(example_noisy,
↳min=-1, max=1)

            prediction = model(example_noisy)
            prediction = torch.reshape(prediction,
↳shape=(128, 1, 28, 28)) # Reconstruct original image dimension
            loss = criterion(prediction, example)

            loss_test += loss
            examples_test.append(example_noisy)
            predictions_test.append(prediction)
        loss_test = loss_test / len(dataloader_test)
        total_loss_test[epoch] = loss_test
        print("Test loss: ", round(loss_test.item(), 4))

    fig, axs = plt.subplots(2, 10, figsize=(20, 5))

    random_nmbrs = np.random.
↳choice(range(len(predictions_test)), size=10)

    # I decided to visualize the noisy image here, as it
↳will show the effect of the noisy autoencoder
    for i in range(10):
        example_image = examples_test[i].cpu().
↳detach()[0][0]

        prediction_image = predictions_test[i].cpu().
↳detach()[0][0]

        fig.suptitle(f'Epoch: {epoch}; Training loss:
↳{loss_train:.4f} Testing loss: {loss_test:.4f}')
        axs[0][i].imshow(example_image, cmap='gray')

```

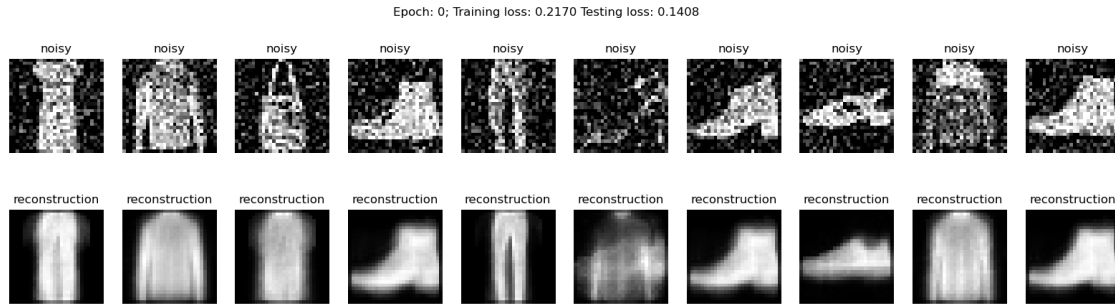
```

        axs[0][i].set_title("noisy")
        axs[0][i].axis('off')
        axs[1][i].imshow(prediction_image, cmap='gray')
        axs[1][i].set_title("reconstruction")
        axs[1][i].axis('off')

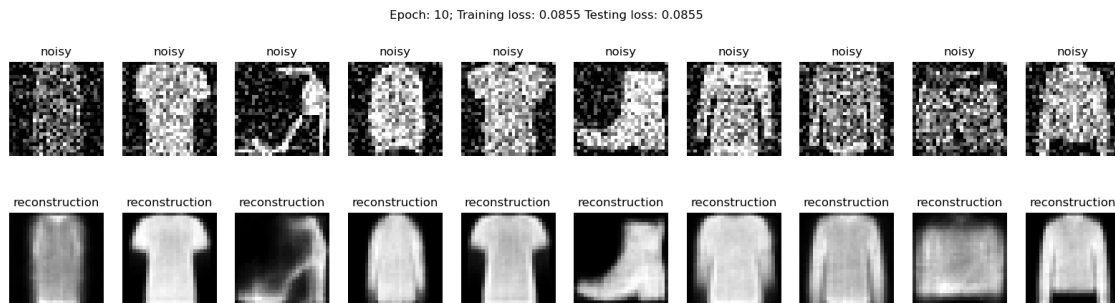
plt.show()

```

Epoch: 0, Training loss: 0.217  
 Test loss: 0.1408

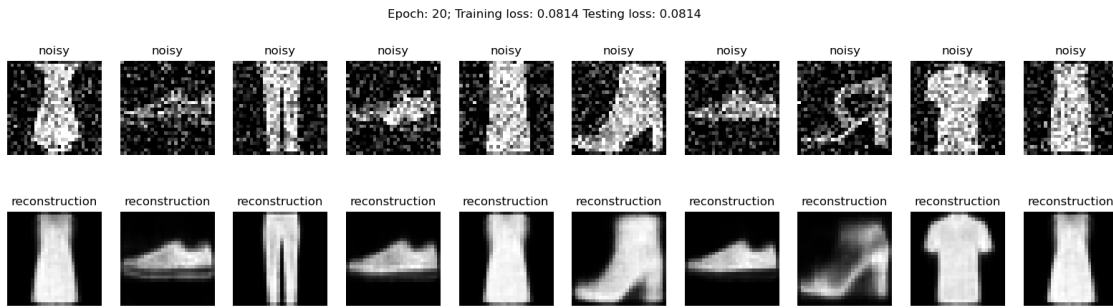


Epoch: 1, Training loss: 0.1219  
 Epoch: 2, Training loss: 0.1051  
 Epoch: 3, Training loss: 0.0976  
 Epoch: 4, Training loss: 0.0944  
 Epoch: 5, Training loss: 0.0919  
 Epoch: 6, Training loss: 0.0899  
 Epoch: 7, Training loss: 0.0884  
 Epoch: 8, Training loss: 0.0872  
 Epoch: 9, Training loss: 0.0863  
 Epoch: 10, Training loss: 0.0855  
 Test loss: 0.0855

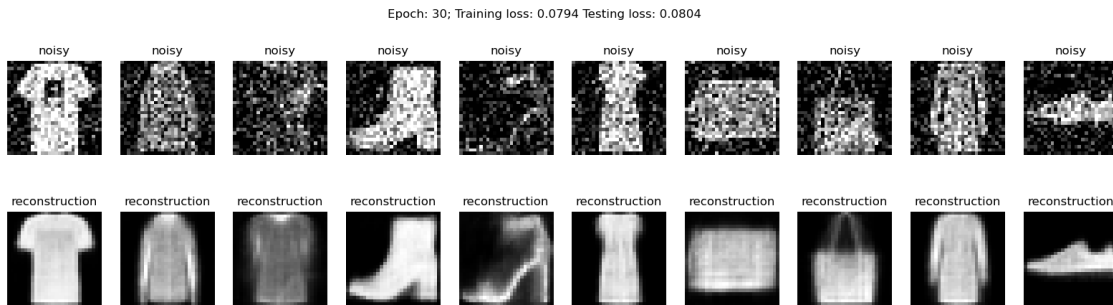


Epoch: 11, Training loss: 0.0848  
 Epoch: 12, Training loss: 0.0842  
 Epoch: 13, Training loss: 0.0837

Epoch: 14, Training loss: 0.0834  
Epoch: 15, Training loss: 0.0829  
Epoch: 16, Training loss: 0.0825  
Epoch: 17, Training loss: 0.0823  
Epoch: 18, Training loss: 0.0819  
Epoch: 19, Training loss: 0.0816  
Epoch: 20, Training loss: 0.0814  
Test loss: 0.0814

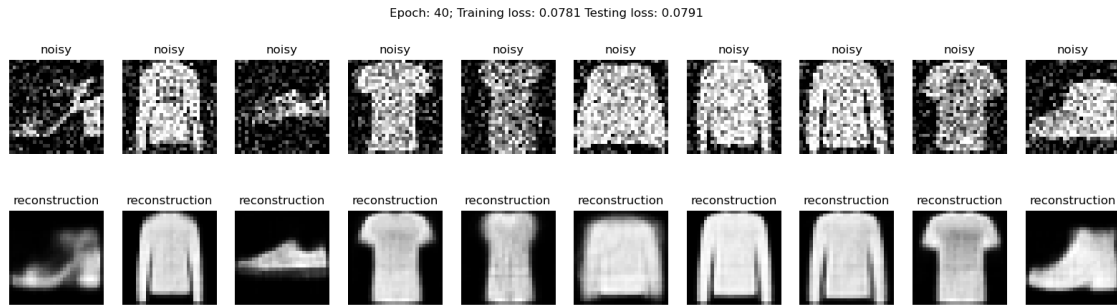


Epoch: 21, Training loss: 0.081  
Epoch: 22, Training loss: 0.0809  
Epoch: 23, Training loss: 0.0807  
Epoch: 24, Training loss: 0.0804  
Epoch: 25, Training loss: 0.0803  
Epoch: 26, Training loss: 0.0801  
Epoch: 27, Training loss: 0.0799  
Epoch: 28, Training loss: 0.0797  
Epoch: 29, Training loss: 0.0795  
Epoch: 30, Training loss: 0.0794  
Test loss: 0.0804

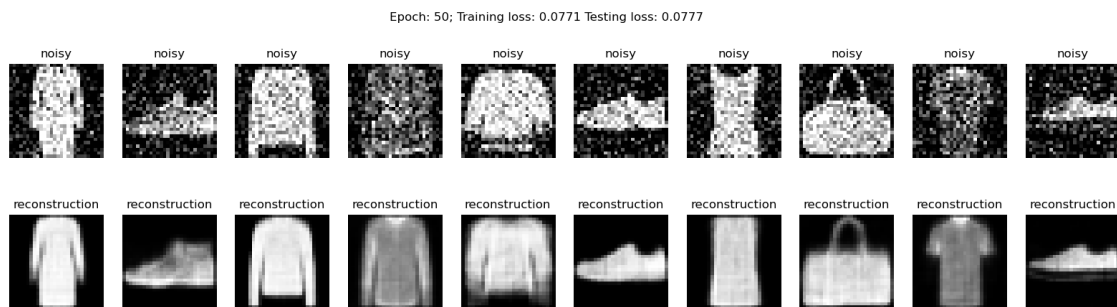


Epoch: 31, Training loss: 0.0792  
Epoch: 32, Training loss: 0.0792  
Epoch: 33, Training loss: 0.079  
Epoch: 34, Training loss: 0.0788

Epoch: 35, Training loss: 0.0787  
 Epoch: 36, Training loss: 0.0786  
 Epoch: 37, Training loss: 0.0783  
 Epoch: 38, Training loss: 0.0783  
 Epoch: 39, Training loss: 0.0781  
 Epoch: 40, Training loss: 0.0781  
 Test loss: 0.0791



Epoch: 41, Training loss: 0.0779  
 Epoch: 42, Training loss: 0.0779  
 Epoch: 43, Training loss: 0.0778  
 Epoch: 44, Training loss: 0.0776  
 Epoch: 45, Training loss: 0.0775  
 Epoch: 46, Training loss: 0.0774  
 Epoch: 47, Training loss: 0.0773  
 Epoch: 48, Training loss: 0.0773  
 Epoch: 49, Training loss: 0.0772  
 Epoch: 50, Training loss: 0.0771  
 Test loss: 0.0777



**Autoencoder** Training loss: 0.0690 Testing loss: 0.0700

**Noisy Autoencoder** Training loss: 0.0771 Testing loss: 0.0777

Both the Training loss and Testing loss are higher when using noisy inputs. This makes sense, since the noise makes it more difficult to reconstruct the original image

```
[ ]: # Save the trained model to be used later
torch.save(model.state_dict(), 'autoencoder_noisy_{}.ckpt'.format(epoch))
```

### 3 Task 3: Implement a Convolutional Autoencoder

```
[ ]: # Define network architecture
class Conv_Autoencoder(nn.Module):
    def __init__(self):
        super(Conv_Autoencoder, self).__init__()
        # Encoder
        self.encoder = nn.Sequential(
            nn.Conv2d(in_channels=1, out_channels=4,
↳kernel_size=5), # 5 x 24 x 24
            nn.ReLU(),
            nn.Conv2d(in_channels=4, out_channels=8,
↳kernel_size=5), # 8 x 20 x 20
            nn.ReLU(),
            nn.Flatten(), # 1 x 3200
            nn.Linear(in_features=3200, out_features=10),
            nn.Softmax()
        )

        # Decoder
        self.decoder = nn.Sequential(
            nn.Linear(in_features=10, out_features=400),
            nn.ReLU(),
            nn.Linear(in_features=400, out_features=4000),
            nn.ReLU(),
            nn.Unflatten(dim=1, unflattened_size=(10, 20, 20)),
            nn.ConvTranspose2d(in_channels=10, out_channels=10,
↳kernel_size=5),
            nn.ReLU(),
            nn.ConvTranspose2d(in_channels=10, out_channels=1,
↳kernel_size=5),
            nn.Tanh()
        )

    def forward(self, x):
        x = x.view(-1, 1, 28, 28)
        x = self.encoder(x)

        x = self.decoder(x)
        x = x.view(-1, 28*28)
        return x

    def generate_image(self, x):
```

```

x = self.decoder(x)
x = x.view(-1, 28*28)
return x

```

### 3.0.1 3.2 Implement reconstruction loss

```
[ ]: criterion = nn.MSELoss()
```

### 3.0.2 3.3 Perform training + 1.4 Visualize results

```
[ ]: # Hyperparameter
lr = 0.001
weight_decay = 1e-5
batch_size = 128
epochs = 51

# Select the device to work on.
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

model = Conv_Autoencoder().to(device)
optimizer = torch.optim.Adam(params=model.parameters(), lr=lr,
    ↪weight_decay=weight_decay)

# Initialize dataloader
dataloader_train = DataLoader(dataset_train, batch_size=batch_size,
    ↪shuffle=True, drop_last=True) # Drop last, to drop last batch since it is
    ↪smaller than 128
dataloader_test = DataLoader(dataset_test, batch_size=batch_size, shuffle=True,
    ↪drop_last=True)

total_loss_train = np.zeros(epochs)
total_loss_test = np.zeros(epochs)
# Perform training
for epoch in range(epochs):
    loss_train = 0
    model.train()
    for [example, _] in dataloader_train:
        example = example.to(device)

        optimizer.zero_grad()
        prediction = model(example)
        prediction = torch.reshape(prediction, shape=(128, 1, 28, 28))
        ↪# Reconstruct original image dimension
        loss = criterion(prediction, example)

    loss_train += loss

```

```

        loss.backward()

        optimizer.step()

    loss_train = loss_train / len(dataloader_train)
    total_loss_train[epoch] = loss_train
    print(f"Epoch: {epoch}, Training loss: {round(loss_train.item(), 4)}")

    # Every 10 epochs, also calculate the loss on the test set
    if epoch % 10 == 0:
        loss_test = 0

        examples_test = []
        predictions_test = []

        model.eval()
        with torch.no_grad():
            for [example, _] in dataloader_test:
                example = example.to(device)

                prediction = model(example)
                prediction = torch.reshape(prediction, ↵
↵shape=(128, 1, 28, 28)) # Reconstruct original image dimension
                loss = criterion(prediction, example)

                loss_test += loss
                examples_test.append(example)
                predictions_test.append(prediction)
            loss_test = loss_test / len(dataloader_test)
            total_loss_test[epoch] = loss_test
            print("Test loss: ", round(loss_test.item(), 4))

        fig, axs = plt.subplots(2, 10, figsize=(20, 5))

        random_nmbrs = np.random.
↵choice(range(len(predictions_test)), size=10)

        for i in range(10):
            example_image = examples_test[i].cpu().
↵detach()[0][0]

            prediction_image = predictions_test[i].cpu().
↵detach()[0][0]

            fig.suptitle(f'Epoch: {epoch}; Training loss: ↵
↵{loss_train:.4f} Testing loss: {loss_test:.4f}')

```



```

        axs[0][i].imshow(example_image, cmap='gray')
        axs[0][i].set_title("original")
        axs[0][i].axis('off')
        axs[1][i].imshow(prediction_image, cmap='gray')
        axs[1][i].set_title("reconstruction")
        axs[1][i].axis('off')

plt.show()

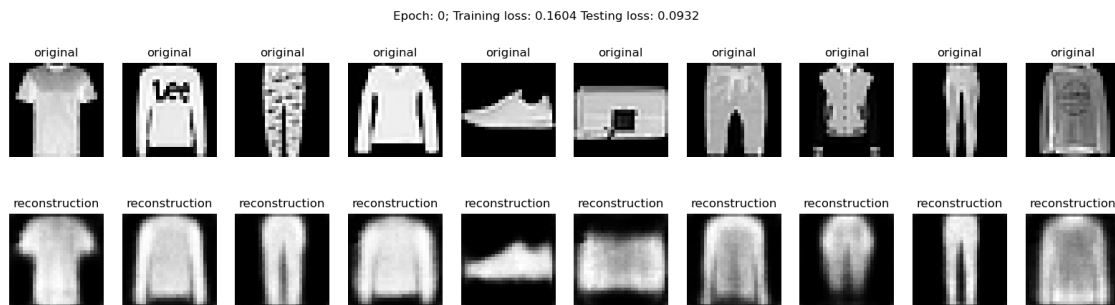
```

c:\Users\adria\anaconda3\envs\genai\Lib\site-packages\torch\nn\modules\module.py:1532: UserWarning: Implicit dimension choice for softmax has been deprecated. Change the call to include dim=X as an argument.

```
return self._call_impl(*args, **kwargs)
```

Epoch: 0, Training loss: 0.1604

Test loss: 0.0932



Epoch: 1, Training loss: 0.084

Epoch: 2, Training loss: 0.0751

Epoch: 3, Training loss: 0.0709

Epoch: 4, Training loss: 0.068

Epoch: 5, Training loss: 0.0659

Epoch: 6, Training loss: 0.064

Epoch: 7, Training loss: 0.0627

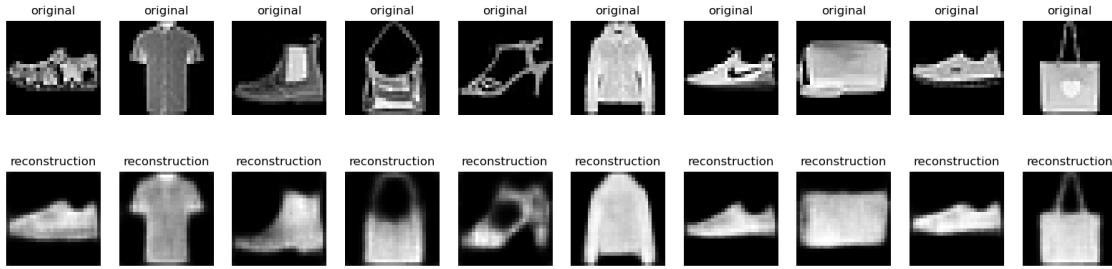
Epoch: 8, Training loss: 0.0615

Epoch: 9, Training loss: 0.0604

Epoch: 10, Training loss: 0.0595


Test loss: 0.0594

Epoch: 10; Training loss: 0.0595 Testing loss: 0.0594



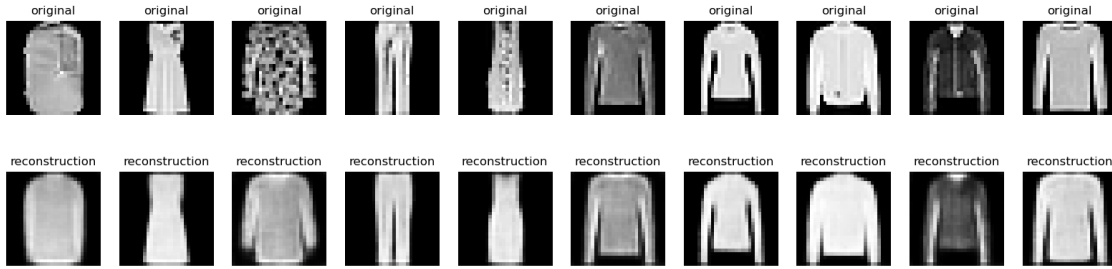
Epoch: 11, Training loss: 0.0588  
Epoch: 12, Training loss: 0.058  
Epoch: 13, Training loss: 0.0573  
Epoch: 14, Training loss: 0.0567  
Epoch: 15, Training loss: 0.0563  
Epoch: 16, Training loss: 0.0558  
Epoch: 17, Training loss: 0.0554  
Epoch: 18, Training loss: 0.0549  
Epoch: 19, Training loss: 0.0546  
Epoch: 20, Training loss: 0.0542  
Test loss: 0.0547

Epoch: 20; Training loss: 0.0542 Testing loss: 0.0547



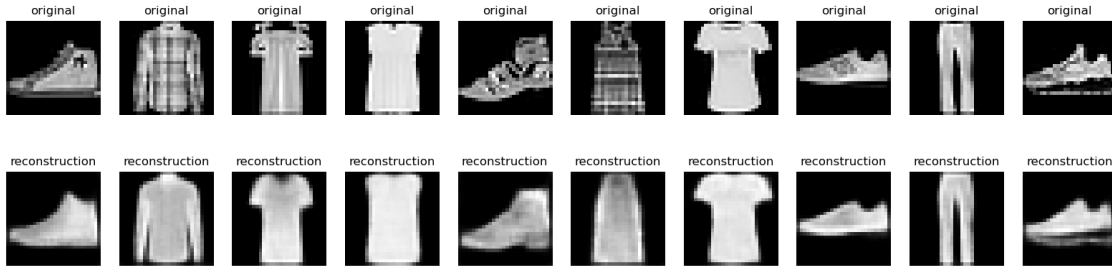
Epoch: 21, Training loss: 0.0539  
Epoch: 22, Training loss: 0.0536  
Epoch: 23, Training loss: 0.0534  
Epoch: 24, Training loss: 0.0532  
Epoch: 25, Training loss: 0.0529  
Epoch: 26, Training loss: 0.0528  
Epoch: 27, Training loss: 0.0525  
Epoch: 28, Training loss: 0.0525  
Epoch: 29, Training loss: 0.0521  
Epoch: 30, Training loss: 0.0521  
Test loss: 0.0532

Epoch: 30; Training loss: 0.0521 Testing loss: 0.0532

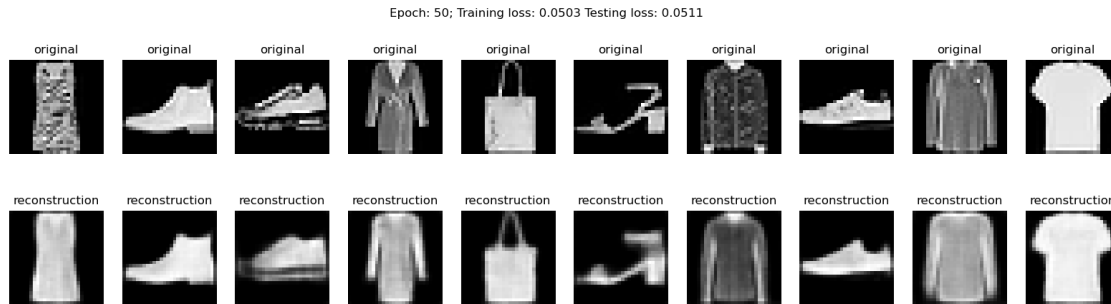


Epoch: 31, Training loss: 0.052  
Epoch: 32, Training loss: 0.0516  
Epoch: 33, Training loss: 0.0517  
Epoch: 34, Training loss: 0.0515  
Epoch: 35, Training loss: 0.0514  
Epoch: 36, Training loss: 0.0513  
Epoch: 37, Training loss: 0.0513  
Epoch: 38, Training loss: 0.0512  
Epoch: 39, Training loss: 0.0509  
Epoch: 40, Training loss: 0.051  
Test loss: 0.0523

Epoch: 40; Training loss: 0.0510 Testing loss: 0.0523



Epoch: 41, Training loss: 0.0509  
Epoch: 42, Training loss: 0.0508  
Epoch: 43, Training loss: 0.0507  
Epoch: 44, Training loss: 0.0506  
Epoch: 45, Training loss: 0.0506  
Epoch: 46, Training loss: 0.0506  
Epoch: 47, Training loss: 0.0504  
Epoch: 48, Training loss: 0.0504  
Epoch: 49, Training loss: 0.0504  
Epoch: 50, Training loss: 0.0503  
Test loss: 0.0511



```
[ ]: # Save the trained model to be used later
torch.save(model.state_dict(), 'autoencoder_conv_{}.ckpt'.format(epoch))
```

### 3.0.3 3.5 Model comparison

**Linear model:** - Training loss: 0.0675 - Test loss: 0.0688

**Convolutional model:** - Training loss: 0.0487 - Test loss: 0.0497

As shown above, both the training loss as well as the test loss are lower for the convolutional model when run with the same hyperparameters. This clearly shows, that the convolutional model is the better model in this case. That makes sense, since convolutional layers work well with image data, since they can detect shapes in local regions.

```
[ ]: # Try out if I can generate an image from arbitrary values. Doesn't really
      ↪work, since we do not know the data distribution in the latent space
rand_x = torch.rand(1,10).to(device)

image = model.generate_image(rand_x)
image = torch.reshape(image, (28, 28))
print(image.shape)
plt.imshow(image.cpu().detach(), cmap='gray')

torch.Size([28, 28])
```

```
[ ]: <matplotlib.image.AxesImage at 0x1b5d21bf950>
```

