

wild-12824010

June 12, 2024

```
[ ]: import torch
import torch.nn as nn

from torchvision.datasets import MNIST, FashionMNIST
from torch.utils.data import DataLoader
from torchvision import transforms
import matplotlib.pyplot as plt
import numpy as np
import torch.nn.functional as F
```

1 Task 1: Training a Variational Autoencoder on MNIST

```
[ ]: class VAE(nn.Module):

    def __init__(self, num_channels=1, num_classes=10, latent_dim=2,
    ↪embed_dim=16):
        super(VAE, self).__init__()
        self.logscale = nn.Parameter(torch.Tensor([0.0])) # Create trainable
    ↪logscale parameter for the p_rec distribution

        self.latent_dim = latent_dim
        self.embedding = nn.Embedding(num_embeddings=num_classes,
    ↪embedding_dim=embed_dim)

        self.encoder = nn.ModuleList([
            nn.Conv2d(in_channels=num_channels, out_channels=8, kernel_size=3,
    ↪stride=2, padding=1),
            nn.Conv2d(in_channels=8, out_channels=16, kernel_size=3, stride=2,
    ↪padding=1),
            nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3, stride=2,
    ↪padding=1),
        ])

        self.decoder = nn.ModuleList([
            nn.Conv2d(in_channels=32, out_channels=16, kernel_size=3,
    ↪padding=1),
```

```

        nn.Conv2d(in_channels=16, out_channels=8, kernel_size=3),
        nn.Conv2d(in_channels=8, out_channels=num_channels, kernel_size=3,
padding=1),
    ])
    self.fc_latent = nn.Linear(in_features=latent_dim + embed_dim,
out_features=512)

    self.fc_mean = nn.Linear(in_features=512 + embed_dim,
out_features=latent_dim)
    self.fc_var = nn.Linear(in_features=512 + embed_dim,
out_features=latent_dim)
    self.leaky_relu = nn.LeakyReLU()
    self.sigmoid = nn.Sigmoid()

def forward(self, x, y):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        y (tensor): Class label(s) of shape [B,].

    Returns:
        x_recon (tensor): Reconstructed image(s) of shape [B, C, H, W].
        mean (tensor): Mean of shape [B, latent_dim].
        log_var (tensor): Log variance of shape [B, latent_dim].
    """
    mean, log_var = self.encode(x, y)
    # Reparameterization Trick
    eps = torch.randn(log_var.shape, device=log_var.device)
    z = mean + torch.exp(log_var * 0.5) * eps
    x_recon = self.decode(z, y)
    return x_recon, mean, log_var

def encode(self, x, y):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        y (tensor): Class label(s) of shape [B,].

    Returns:
        mean (tensor): Mean of shape [B, latent_dim].
        log_var (tensor): Log variance of shape [B, latent_dim].
    """
    for layer in self.encoder:
        x = layer(x)
        x = self.leaky_relu(x)
    x = torch.reshape(x, (x.shape[0], -1))
    class_embed = self.embedding(y)

```

```

    # Concat class information
    mean = self.fc_mean(torch.cat((x, class_embed), dim=1))
    log_var = self.fc_var(torch.cat((x, class_embed), dim=1))
    return mean, log_var

def decode(self, z, y):
    """
    Args:
        z (tensor): Latent variable(s) of shape [B, latent_dim].
        y (tensor): Class label(s) of shape [B,].

    Returns:
        x (tensor): Reconstructed image(s) of shape [B, C, H, W].
    """
    class_embed = self.embedding(y)
    # Concat class information
    x = self.fc_latent(torch.cat((z, class_embed), dim=1))
    x = torch.reshape(x, (-1, 32, 4, 4))
    for layer in self.decoder:
        x = nn.functional.interpolate(x, scale_factor=2, mode='bilinear',
↪align_corners=True)
        x = self.leaky_relu(x)
        x = layer(x)
    x = self.sigmoid(x)
    return x

def sample(self, y, device):
    """
    Args:
        y (int): Class label.
        device (torch.device): Which device to use (cuda or cpu).

    Returns:
        (tensor): Image of shape [1, C, H, W].
    """
    z = torch.randn((1, self.latent_dim), device=device)
    return self.decode(z, torch.tensor([y], device=device))

def sample_latent(self, x, y):
    """
    Args:
        x (tensor): Image(s) of shape [B, C, H, W].
        y (tensor): Class label(s) of shape [B,].

    Returns:
        z (tensor): Latent variable(s) of shape [B, latent_dim].
    """

```

```

mean, log_var = self.encode(x, y)
# Reparameterization Trick
eps = torch.randn(log_var.shape, device=log_var.device)
z = mean + torch.exp(log_var * 0.5) * eps
return z

```

1.0.1 1.1 Implement ELBO loss function and training routine

```

[ ]: def gaussian_likelihood(x_recon, x, logscale):
    # Calculate the gaussian likelihood of x und x_recon distribution (use it as
    ↪reconstruction loss)
    scale = torch.exp(logscale)
    mean = x_recon
    p_rec = torch.distributions.Normal(mean, scale)
    # measure prob of seeing image under p(x/z)
    log_pxz = p_rec.log_prob(x)
    return -torch.mean(log_pxz.sum(dim=(1, 2, 3)))

def kl_divergence(z, mean, std):
    # Monte carlo KL divergence
    # 1. define the first two probabilities (Normal for both cases here)
    p = torch.distributions.Normal(torch.zeros_like(mean), torch.ones_like(std))
    ↪# fix distribution to N(0,1)
    q = torch.distributions.Normal(mean, std) # Over time, the q distribution
    ↪will move closer to the p distribution

    # 2. get the probabilities from the equation
    log_qzx = q.log_prob(z)
    log_pz = p.log_prob(z)

    # Calculate kl
    kl = (log_qzx - log_pz)
    kl = kl.sum(-1)
    return torch.mean(kl)

def elbo_loss(x_recon, x, mean, log_var, kl_weight, logscale):

    # Sample z from q(z|x)
    z = sample_z(mean, log_var)

    std = torch.exp(log_var / 2)
    kl = kl_divergence(z, mean, std) # kl needs std and not log_var, because it
    ↪can only take pos values
    recon_loss = gaussian_likelihood(x_recon, x, logscale)

    loss = (kl * kl_weight + recon_loss)
    return loss

```

```
def sample_z(mean, log_var):
    # sample z from q(z|x)
    std = torch.exp(log_var / 2)
    q = torch.distributions.Normal(mean, std)
    z = q.rsample()
    return z
```

```
[ ]: # Alternative loss function which I am not using
def vae_loss(x_recon, x, mean, log_var, kl_weight=0.001):
    recon_loss = F.binary_cross_entropy(x_recon, x, reduction='sum')
    kl_div = -0.5 * torch.sum(1 + log_var - mean.pow(2) - log_var.exp())
    batch_size = x.size(0)
    return (recon_loss + kl_weight * kl_div) / batch_size

def reconstruction_loss(x_recon, x):
    loss = torch.sum(torch.square(x_recon-x))
    return torch.mean(loss)

def kl_loss(mean, log_var):
    loss = -0.5 * torch.sum(1 + log_var - torch.square(mean) - torch.
↪square(torch.exp(log_var)), axis=-1)
    return torch.mean(loss)
```

```
[ ]: def plot_samples(model, device):
    model.eval()
    with torch.no_grad():
        fig, axs = plt.subplots(1, 10, figsize=(10, 1))
        for i in range(10):
            sample = model.sample(i, device).cpu().squeeze()
            axs[i].set_title(i)
            axs[i].imshow(sample, cmap='gray')
            axs[i].axis('off')
    plt.show()
```

```
[ ]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')

transform = transforms.Compose([
    transforms.ToTensor()
])
dataset_train = MNIST('./data', train=True, download=True, transform=transform)
dataset_test = MNIST('./data', train=False, download=True, transform=transform)
```

```
[ ]: # Hyperparameters
num_epochs = 5
```

```

kl_weight = 1 # For me, setting the weight to 0.0001 heavily harms the
↳ performance of the model, so I leave it like this
batch_size = 128

train_loader = DataLoader(dataset_train, batch_size=batch_size, shuffle=True)

def training_routine(num_epochs=5, kl_weight=1, loss_function=""):

    model = VAE().to(device)
    optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

    losses = []
    for epoch in range(num_epochs):
        model.train()
        train_loss = 0
        num_batches = 0

        for x, y in train_loader:
            x, y = x.to(device), y.to(device)
            optimizer.zero_grad()
            x_recon, mean, log_var = model(x, y)
            if loss_function == "recon":
                loss = gaussian_likelihood(x_recon, x, model.logscale) / x.shape[0] #
↳ scale loss down
            elif loss_function == "kl":
                # sample z from q(z/x)
                std = torch.exp(log_var / 2)
                q = torch.distributions.Normal(mean, std)
                z = q.rsample()
                loss = kl_divergence(z, mean, std)
            else:
                loss = elbo_loss(x_recon, x, mean, log_var, kl_weight, model.logscale) /
↳ x.shape[0] # scale loss down

            loss.backward()
            train_loss += loss.item() * x.shape[0]
            num_batches += x.shape[0]
            optimizer.step()
            losses.append(loss.cpu().detach().numpy())

        avg_train_loss = train_loss / num_batches # Divide by sample number

    print(f'Epoch [{epoch+1}/{num_epochs}], Loss: {avg_train_loss:.4f}')


    plot_samples(model, device)

```


```
plt.plot(losses)
plt.xlabel('Step')
plt.ylabel('Loss')
plt.title('Training Curve')
plt.show()
return model
```

```
[ ]: model = training_routine()
```


Epoch [1/5], Loss: 4.4815




Epoch [2/5], Loss: 1.8550





Epoch [3/5], Loss: -0.1203



Epoch [4/5], Loss: -1.2014



Epoch [5/5], Loss: -1.5260



2 Task 2: Visualize the latent space

```
[ ]: def visualize_latent_space(model, loss_name ):
    latent_variables = []
    labels = []
    dataloader_test = DataLoader(dataset_test, batch_size=1, shuffle=True)

    for x, y in dataloader_test:
        labels.append(y[0].detach().numpy().flatten().flatten())
        x = x.to(device)
        y = y.to(device)
        z = model.sample_latent(x, y)
```



```

latent_variables.append(z.cpu().detach().numpy().flatten().flatten())
labels = [label[0] for label in labels]

# Visualize latent space
latent_variables = np.array(latent_variables)
unique_labels = np.unique(labels)

plt.figure(figsize=(10, 7))

# Create a scatter plot for each unique label
for label in unique_labels:
    mask = labels == label
    plt.scatter(latent_variables[mask, 0], latent_variables[mask, 1],
        label=label, alpha=0.5, s=10)
plt.title(f"Latent space visualization for {loss_name}")
plt.legend(title='Numbers')
plt.show()


```

```


[ ]: elbo_model = training_routine()
visualize_latent_space(elbo_model, "ELBO loss (combined loss) 1")

```


Epoch [1/5], Loss: 4.4801




Epoch [2/5], Loss: 1.8863





Epoch [3/5], Loss: -0.0908

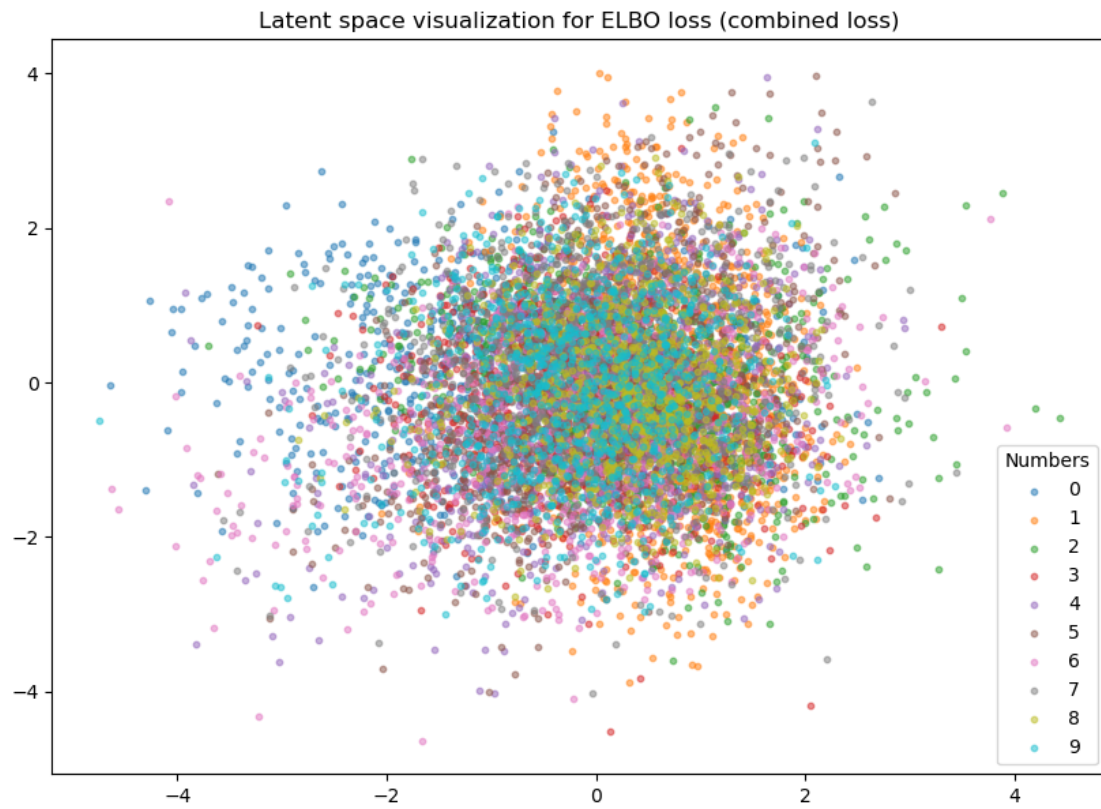


Epoch [4/5], Loss: -1.1789




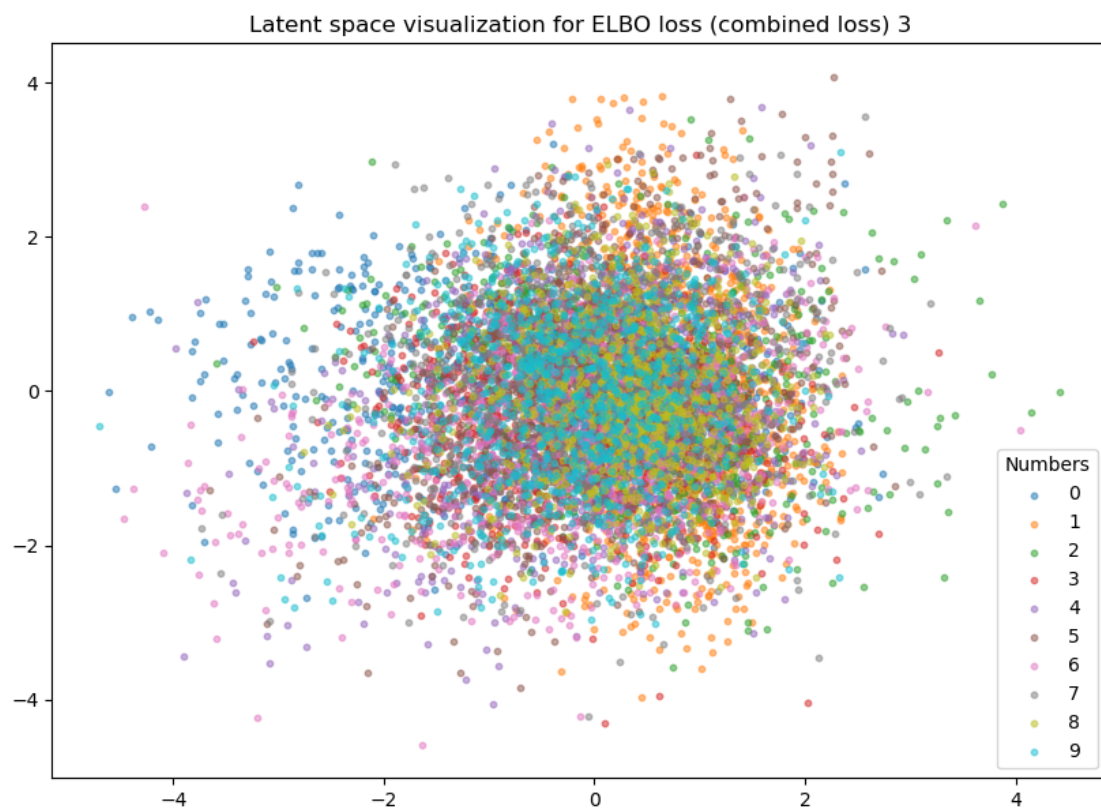
Epoch [5/5], Loss: -1.5082






```
[ ]: visualize_latent_space(elbo_model, "ELBO loss (combined loss) 2")  
visualize_latent_space(elbo_model, "ELBO loss (combined loss) 3")
```






```
[ ]: recon_model = training_routine(loss_function="recon")
visualize_latent_space(recon_model, "Reconstruction loss")
```


Epoch [1/5], Loss: 4.4688




Epoch [2/5], Loss: 1.8482




Epoch [3/5], Loss: -0.1192




Epoch [4/5], Loss: -1.1899




Epoch [5/5], Loss: -1.5142









```
[ ]: visualize_latent_space(recon_model, "Reconstruction loss 2")  
visualize_latent_space(recon_model, "Reconstruction loss 3")
```




```
[ ]: kl_model = training_routine(loss_function="kl")  
visualize_latent_space(kl_model, "KL divergence")
```


Epoch [1/5], Loss: 0.0017




Epoch [2/5], Loss: 0.0012




Epoch [3/5], Loss: 0.0012





Epoch [4/5], Loss: 0.0010




Epoch [5/5], Loss: 0.0011








```
[ ]: visualize_latent_space(kl_model, "KL divergence 2")  
visualize_latent_space(kl_model, "KL divergence 3")
```






3 Task 3: Anomaly Detection using a Variational Autoencoder


```
[ ]: model_task3 = training_routine(num_epochs=5)
```

```
c:\Users\adria\anaconda3\envs\genai\Lib\site-
packages\torch\autograd\graph.py:744: UserWarning: Plan failed with a
cudnnException: CUDNN_BACKEND_EXECUTION_PLAN_DESCRIPTOR: cudnnFinalize
Descriptor Failed cudnn_status: CUDNN_STATUS_NOT_SUPPORTED (Triggered internally
at C:\cb\pytorch_1000000000000\work\aten\src\ATen\native\cudnn\Conv_v8.cpp:919.)
    return Variable._execution_engine.run_backward( # Calls into the C++ engine
to run the backward pass
```


Epoch [1/5], Loss: 4.4770




Epoch [2/5], Loss: 1.8787




Epoch [3/5], Loss: -0.0822




Epoch [4/5], Loss: -1.1610



Epoch [5/5], Loss: -1.4877





```
[ ]: fashion_dataset_train = FashionMNIST('./data', train=True, download=True,
    ↪transform=transform)

[ ]: dataloader_train3_fashion = DataLoader(fashion_dataset_train, batch_size=1,
    ↪shuffle=True)
dataloader_train3 = DataLoader(dataset_train, batch_size=1, shuffle=True)

model_task3.eval()
with torch.no_grad():

    fig, axs = plt.subplots(1,10, figsize=(20,3))

    for i, sample in enumerate(dataloader_train3):
        x = sample[0].to(device)
        y = sample[1].to(device)
        x_recon, mean, log_var = model_task3.forward(x, y)
        loss = elbo_loss(x_recon, x, mean, log_var, kl_weight, model_task3.logscale)

        axs[i].imshow(sample[0].squeeze(), cmap='gray')
```

```

    axs[i].axis('off')
    axs[i].set_title(f"Loss: {loss:.2}")
    if i == 9:
        break
fig.suptitle("Reconstruction using MNIST dataset")
fig.show()

fig, axs = plt.subplots(1, 10, figsize=(20,3))
for i, sample in enumerate(dataloader_train3_fashion):
    x = sample[0].to(device)
    y = sample[1].to(device)
    x_recon, mean, log_var = model_task3.forward(x, y)
    loss = elbo_loss(x_recon, x, mean, log_var, kl_weight, model_task3.logscale)

    axs[i].imshow(sample[0].squeeze(), cmap='gray')
    axs[i].axis('off')
    axs[i].set_title(f"Loss: {loss:.2}")
    if i == 9:
        break
fig.suptitle("Reconstruction using MNISTFashion dataset")
fig.show()

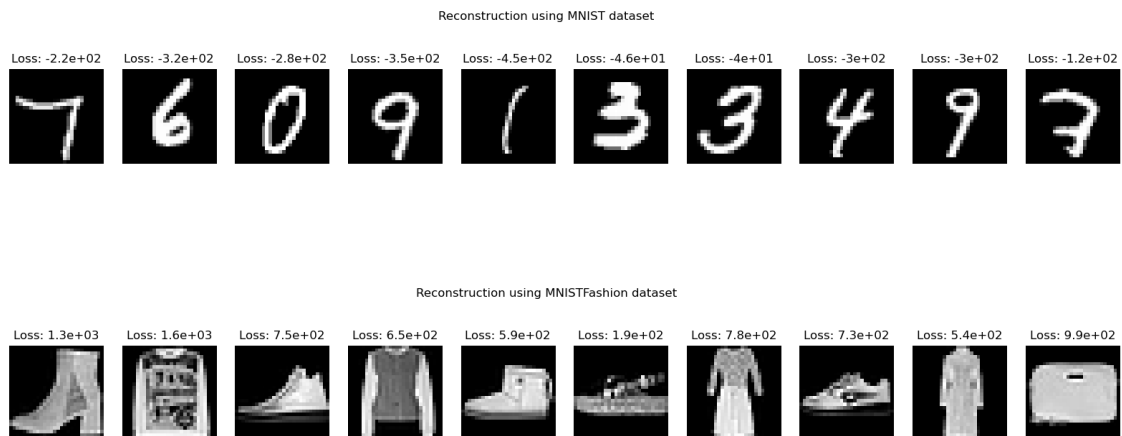
```

C:\Users\adria\AppData\Local\Temp\ipykernel_8136\2371798454.py:22: UserWarning: FigureCanvasAgg is non-interactive, and thus cannot be shown

fig.show()

C:\Users\adria\AppData\Local\Temp\ipykernel_8136\2371798454.py:38: UserWarning: FigureCanvasAgg is non-interactive, and thus cannot be shown

fig.show()



The loss is usually way higher when using the MNISTFashion dataset. Thus, this model can be used for anomaly detection

```
[ ]: sample = model_task3.sample(torch.tensor(9), device).cpu().detach().numpy()  
plt.imshow(sample.squeeze(), cmap='gray')
```

```
[ ]: <matplotlib.image.AxesImage at 0x20ac2833500>
```

